

# Concurrent Robin Hood Hashing<sup>1</sup>

Supervisor: Barak A. Pearlmutter<sup>2</sup>



# Motivations<sup>1</sup>

- Make improvements on a 10 year old state of the art.
- Provide the first concurrent Robin Hood Hashing in the literature.

# Contributions<sup>3</sup>

- First linearisable concurrent variant of Robin Hood Hashing.
- Strong application of new *K-CAS* developments [Arbel-Raviv,, Brown; 2016]
- Competitive performance compared to state of the art concurrent hash tables.

# General talk structure<sup>1</sup>

- Hash table and Robin Hood background<sup>2</sup>
- Challenges with concurrent Robin Hood
- What are the options?
- Solution
- Correctness/Progress
- Evaluation

# Hash Tables<sup>1</sup>

- Constant time  $O(1)$  set/map structures<sup>2</sup>
- Set operations:

1. **Contains(Key)**<sup>3</sup>
2. **Add(Key)**
3. **Remove(Key)**

- No need for sorting of keys, unlike tree-based sets/maps<sup>4</sup>
- Require a hash function for keys
- Applications: Search, Object representation in VMs/interpreters, caches...

# Hash Tables<sup>1</sup>

Divided into two camps: Open vs Closed Addressing.<sup>2</sup>

## Open Addressing.<sup>3</sup>

- Items are stored in individual buckets only.
- If bucket is already taken find a new one: Collision algorithm.

## Closed Addressing.<sup>5</sup>

- Items are stored at original bucket only.
- Typically in a linked list structure.



# Robin Hood Hashing<sup>2</sup> (Open Addressing)

# Robin Hood [Celis ;86]<sup>1</sup>

# Robin Hood [Celis ;86]<sup>1</sup>

Motto: Steal from the rich<sup>2</sup>  
and give to the poor.



# Robin Hood [Celis ;86]<sup>1</sup>

Motto: Steal from the rich<sup>2</sup>  
and give to the poor.

Search: Linear probing with<sup>3</sup>  
culling.

# Robin Hood [Celis ;86]<sup>1</sup>

Motto: Steal from the rich<sup>2</sup>  
and give to the poor.

Search: Linear probing with<sup>3</sup>  
culling.

Insertion: Linear probing with<sup>4</sup>  
conditional recursive  
displacement.

# Robin Hood [Celis ;86]<sup>1</sup>

Motto: Steal from the rich<sup>2</sup>  
and give to the poor.

Search: Linear probing with<sup>3</sup>  
culling.

Insertion: Linear probing with<sup>4</sup>  
conditional recursive  
displacement.

Removal: Backward shifting.<sup>5</sup>  
More on that later.

# Robin Hood [Celis ;86]<sup>1</sup>

Motto: Steal from the rich<sup>2</sup>  
and give to the poor.

Search: Linear probing with<sup>3</sup>  
culling.

Insertion: Linear probing with<sup>5</sup>  
conditional recursive  
displacement.

Removal: Backward shifting.<sup>6</sup>  
More on that later.

Definition: The number of buckets<sup>4</sup>  
away an entry is from its ideal  
bucket - **D**istance **F**rom **B**ucket  
(*DFB*)

# Robin Hood [Celis ;86]<sup>1</sup>

Motto: Steal from the rich and give to the poor.<sup>2</sup>

Search: Linear probing with culling.<sup>3</sup>

Insertion: Linear probing with conditional recursive displacement.<sup>5</sup>

Removal: Backward shifting. More on that later.<sup>7</sup>

Definition: The number of buckets away an entry is from its ideal bucket - **D**istance **F**rom **B**ucket (*DFB*)<sup>4</sup>

If relocated item has bigger *DFB* than than current, kick current out, take spot, and recursively insert current further down the table.<sup>6</sup>

# Linear Probing vs Robin Hood<sup>1</sup>

Initial Table, inserting V.<sup>2</sup>

Ø	X <sub>0</sub>	Y <sub>1</sub>	Z <sub>1</sub>	W <sub>1</sub>	Ø
---	----------------	----------------	----------------	----------------	---

<sup>3</sup>

# Linear Probing vs Robin Hood<sup>1</sup>


Initial Table, inserting V.<sup>2</sup>

Key:<sup>3</sup>

<sup>5</sup>- Moved item
- Inserted item

<sup>4</sup>  

Ø	X <sub>0</sub>	Y <sub>1</sub>	Z <sub>1</sub>	W <sub>1</sub>	Ø
---	----------------	----------------	----------------	----------------	---



Linear Probing Table<sup>7</sup>

<sup>8</sup>  

Ø	X <sub>0</sub>	Y <sub>1</sub>	Z <sub>1</sub>	W <sub>1</sub>	V <sub>4</sub>
---	----------------	----------------	----------------	----------------	----------------

# Linear Probing vs Robin Hood


Initial Table, inserting V.<sup>2</sup>

Key:<sup>3</sup>

<sup>4</sup>- Moved item
- Inserted item

<sup>5</sup>  

Ø	X <sub>0</sub>	Y <sub>1</sub>	Z <sub>1</sub>	W <sub>1</sub>	Ø
---	----------------	----------------	----------------	----------------	---



Linear Probing Table<sup>7</sup>

<sup>8</sup>  

Ø	X <sub>0</sub>	Y <sub>1</sub>	Z <sub>1</sub>	W <sub>1</sub>	V <sub>4</sub>
---	----------------	----------------	----------------	----------------	----------------

Robin Hood Table<sup>9</sup>

<sup>10</sup>  

Ø	X <sub>0</sub>	Y <sub>1</sub>	V <sub>2</sub>	Z <sub>2</sub>	W <sub>2</sub>
---	----------------	----------------	----------------	----------------	----------------




# Linear Probing vs Robin Hood

Initial Table, inserting V.

Key:

- Moved item
- Inserted item

Ø	X <sub>0</sub>	Y <sub>1</sub>	Z <sub>1</sub>	W <sub>1</sub>	Ø
---	----------------	----------------	----------------	----------------	---



Linear Probing Table

Ø	X <sub>0</sub>	Y <sub>1</sub>	Z <sub>1</sub>	W <sub>1</sub>	V <sub>4</sub>
---	----------------	----------------	----------------	----------------	----------------

- Less work

Robin Hood Table

Ø	X <sub>0</sub>	Y <sub>1</sub>	V <sub>2</sub>	Z <sub>2</sub>	W <sub>2</sub>
---	----------------	----------------	----------------	----------------	----------------

- Less distance variance

# Robin Hood Search<sup>1</sup>

# Robin Hood Search<sup>1</sup>

Linear probe as normal.<sup>2</sup>

# Robin Hood Search<sup>1</sup>

Linear probe as normal.<sup>2</sup>

When you see someone not  
as far away as you: Stop.<sup>3</sup>

# Robin Hood Search<sup>1</sup>

Linear probe as normal.<sup>2</sup>


$U_0$					
$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$

When you see someone not  
as far away as you: Stop.<sup>4</sup>

# Robin Hood Search<sup>1</sup>

Linear probe as normal.<sup>2</sup>

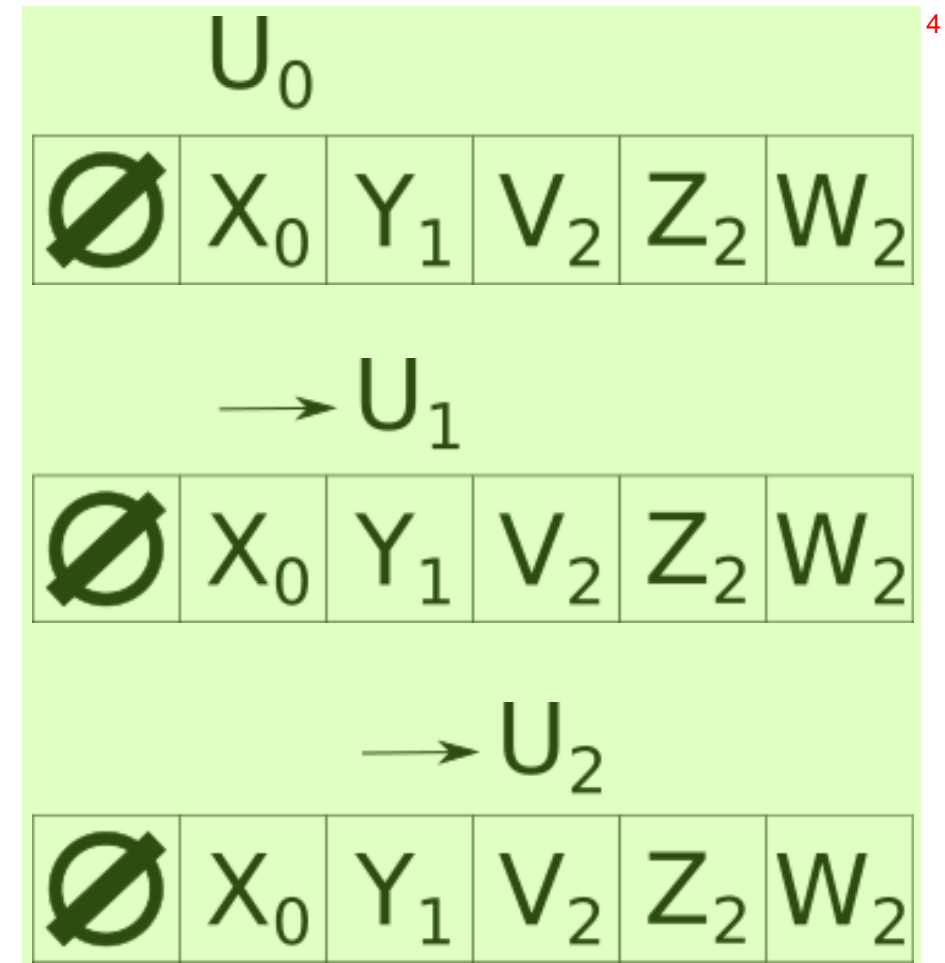
When you see someone not  
as far away as you: Stop.<sup>3</sup>



# Robin Hood Search<sup>1</sup>

Linear probe as normal.<sup>2</sup>

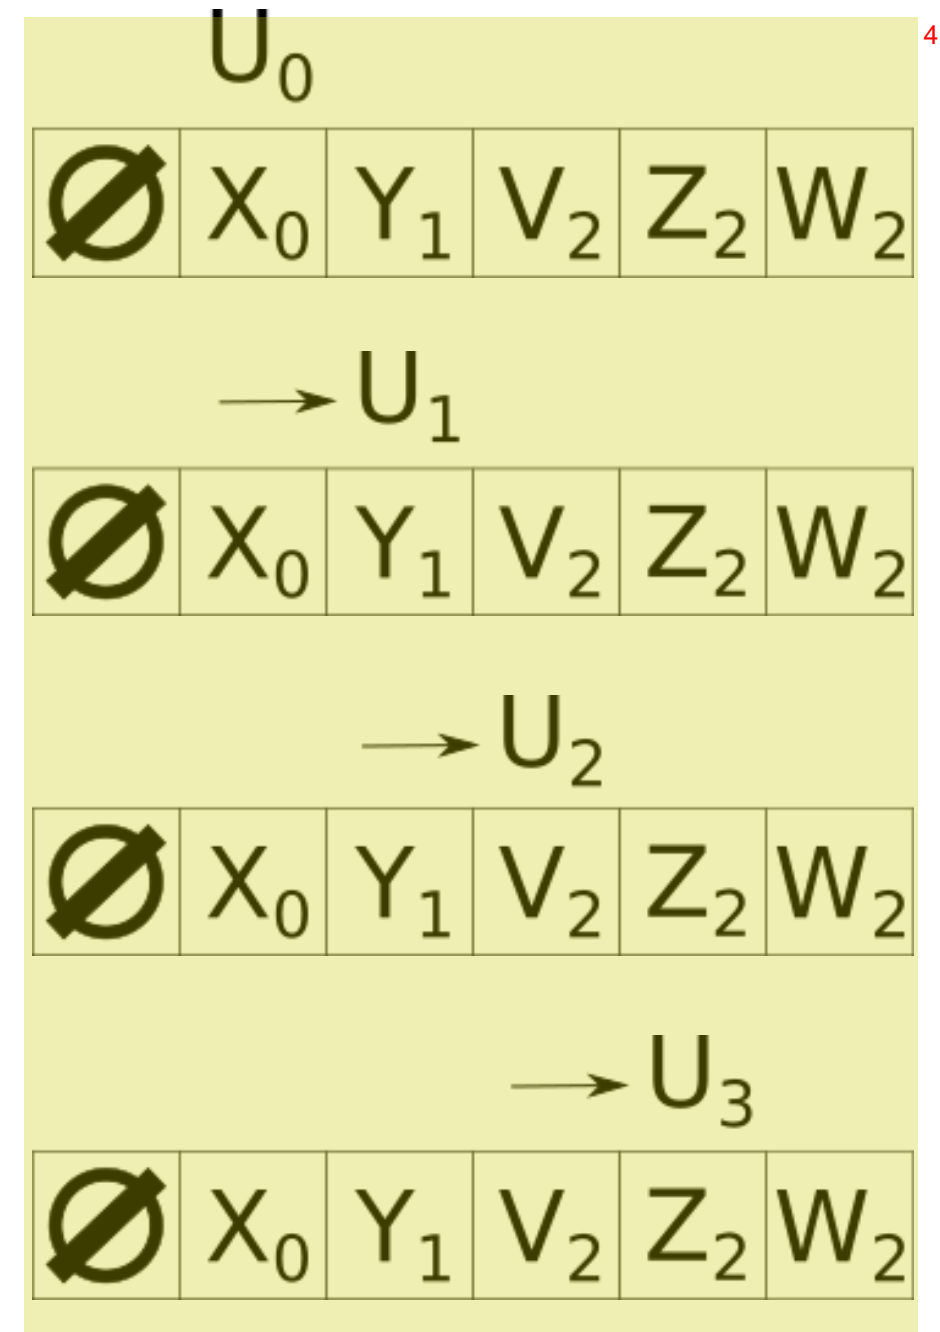
When you see someone not as far away as you: Stop.<sup>3</sup>



# Robin Hood Search<sup>1</sup>

Linear probe as normal.<sup>2</sup>

When you see someone not as far away as you: Stop.<sup>3</sup>





# Robin Hood benefits<sup>1</sup>

# Robin Hood benefits<sup>1</sup>

## 1. Fast, predictable performance:<sup>2</sup>

- Optimised for reads - 2.6 probes per successful search.<sup>3</sup>
- $\log(n)$  on failed search.
- Doesn't degenerate over time (poisoning).

# Robin Hood benefits<sup>1</sup>

## 1. Fast, predictable performance:<sup>2</sup>

- Optimised for reads - 2.6 probes per successful search.<sup>3</sup>
- $\log(n)$  on failed search.
- Doesn't degenerate over time (poisoning).

## 2. Relatively simple:<sup>4</sup>

- No linked list or pointer manipulation.<sup>5</sup>

# Robin Hood benefits<sup>1</sup>

## 1. Fast, predictable performance:<sup>2</sup>

- Optimised for reads - 2.6 probes per successful search.<sup>3</sup>
- $\log(n)$  on failed search.
- Doesn't degenerate over time (poisoning).


## 2. Relatively simple:<sup>4</sup>

- No linked list or pointer manipulation.<sup>5</sup>


## 3. Cache efficient.<sup>6</sup>

- Flat data, low probes.<sup>7</sup>
- No dynamic allocation.
- Probes are generally on a single cache line.

# Standard Solution: *Sharded* locks




# Standard Solution: *Sharded* locks



- Could grab multiple locks.
- Could result in deadlock, if allowed to wrap around.

# Standard Solution: *Sharded locks*



- Could grab multiple locks.
- Could result in deadlock, if allowed to wrap around.
- Not very clean for our case: Need extra phantom segment to stop deadlock.
- Hacky. Slow. Lots of contention.

# Contention: Remove<sup>1</sup>



# Contention: Remove<sup>1</sup>

First part is normal search.<sup>2</sup>

# Contention: Remove<sup>1</sup>

First part is normal search.<sup>2</sup>

When you find the entry<sup>3</sup>  
delete it, move everyone  
back by 1.

# Contention: Remove<sup>1</sup>

First part is normal search.<sup>2</sup>

When you find the entry<sup>3</sup>  
delete it, move everyone  
back by 1.

Stopping condition: The<sup>4</sup>  
entry you're moving back  
has a distance of 0 or it's a  
Null.

# Contention: Remove<sup>1</sup>

First part is normal search.<sup>2</sup>

When you find the entry delete it, move everyone back by 1.<sup>3</sup>

Stopping condition: The entry you're moving back has a distance of 0 or it's a Null.<sup>4</sup>

Initial table.<sup>5</sup>

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

We want to delete Y.<sup>7</sup>

# Contention: Remove<sup>1</sup>

First part is normal search.<sup>2</sup>

When you find the entry delete it, move everyone back by 1.<sup>3</sup>

Stopping condition: The entry you're moving back has a distance of 0 or it's a Null.<sup>4</sup>

Initial table.<sup>5</sup>

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

We want to delete Y.<sup>7</sup>

Y deleted.<sup>8</sup>

$\emptyset$	$X_0$	$V_1$	$Z_1$	$W_1$	$\emptyset$
-------------	-------	-------	-------	-------	-------------

# Contention: Remove<sup>1</sup>


First part is normal search.<sup>2</sup>

When you find the entry delete it, move everyone back by 1.<sup>3</sup>

Stopping condition: The entry you're moving back has a distance of 0 or it's a Null.<sup>4</sup>


Great source of contention.<sup>5</sup>

Initial table.<sup>6</sup>



We want to delete Y.<sup>8</sup>

Y deleted.<sup>9</sup>



# Problems with creating Concurrent Robin Hood

# Problems with creating Concurrent Robin Hood<sup>1</sup>

- Lots of moving parts. Performance problem.<sup>2</sup>



# Problems with creating Concurrent Robin Hood<sup>1</sup>

- Lots of moving parts. Performance problem.<sup>2</sup>
- An insert can trigger a global table reorganisation.

# Problems with creating Concurrent Robin Hood<sup>1</sup>

- Lots of moving parts. Performance problem.
  - An insert can trigger a global table reorganisation.
  - Cyclic lock grabbing. If locks are *sharded*, deadlock is possible.
  - Huge contention on **Remove**.
- <sup>2</sup>

# Problems with creating Concurrent Robin Hood<sup>1</sup>

- Lots of moving parts. Performance problem.
- An insert can trigger a global table reorganisation.
- Cyclic lock grabbing. If locks are *sharded*, deadlock is possible.
- Huge contention on **Remove**.

## Possible Solutions<sup>3</sup>

- Bespoke non-blocking solution<sup>4</sup>
- Transactional Memory
- *K-CAS* (Multi-word compare and swap)

# Issues with bespoke

# Issues with bespoke<sup>1</sup>

1. Use of dynamic memory, ruins cache locality. Slow.<sup>2</sup>

# Issues with bespoke<sup>1</sup>

1. Use of dynamic memory, ruins cache locality. Slow.<sup>2</sup>

2. Horrendously complicated. Difficult to even to get insertion working. Robin Hood invariant must hold all the time.<sup>3</sup>

# Issues with bespoke<sup>1</sup>

1. Use of dynamic memory, ruins cache locality. Slow.<sup>2</sup>

2. Horrendously complicated. Difficult to even to get insertion working. Robin Hood invariant must hold all the time.<sup>3</sup>

3. End up reimplementing *K*-CAS. Why not just use *K*-CAS?<sup>4</sup>

# Issues with bespoke<sup>1</sup>

1. Use of dynamic memory, ruins cache locality. Slow.<sup>2</sup>

2. Horrendously complicated. Difficult to even to get insertion working. Robin Hood invariant must hold all the time.<sup>3</sup>

3. End up reimplementing *K*-CAS. Why not just use *K*-CAS?<sup>4</sup>


## Method Chosen: *K*-CAS<sup>5</sup>

*K*-CAS is a *multi-word compare-and-swap* primitive. Each table operation is described as one large *K*-CAS.<sup>6</sup>



# Our Solution:<sup>1</sup>

## *Sharded timestamps*<sup>2</sup>



Similar to lock-base *sharding*. Groups of timestamps *protect* the table.<sup>4</sup>

Each relocation operation increments the timestamp. Except relocations can be done in bulk.<sup>5</sup>

# *K-CAS* without timestamps<sup>1</sup>

Say our *K-CAS* solution just encapsulates every modifying operation {**Add**, **Remove**} into a *K-CAS* operation.<sup>2</sup>

# *K-CAS* without timestamps<sup>1</sup>

Say our *K-CAS* solution just encapsulates every modifying operation {**Add**, **Remove**} into a *K-CAS* operation.<sup>2</sup>

Readers of the table may miss entries that are moved backwards by remove.<sup>3</sup>

# *K-CAS* without timestamps<sup>1</sup>

Try find  $V$ .<sup>4</sup>

$V_0$ <sup>5</sup>

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

<sup>6</sup>

Say our *K-CAS* solution just encapsulates every modifying operation {**Add**, **Remove**} into a *K-CAS* operation.<sup>2</sup>

Readers of the table may miss entries that are moved backwards by remove.<sup>3</sup>

# *K-CAS* without timestamps<sup>1</sup>

Say our *K-CAS* solution just encapsulates every modifying operation {**Add**, **Remove**} into a *K-CAS* operation.<sup>2</sup>

Readers of the table may miss entries that are moved backwards by remove.<sup>3</sup>

Try find  $V$ .<sup>4</sup>

$V_0$ <sup>5</sup>

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

<sup>6</sup>

$\rightarrow V_1$ <sup>7</sup>

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

<sup>8</sup>

# *K-CAS* without timestamps<sup>1</sup>

Say our *K-CAS* solution just encapsulates every modifying operation {**Add**, **Remove**} into a *K-CAS* operation.<sup>2</sup>

Readers of the table may miss entries that are moved backwards by remove.<sup>3</sup>

Try find  $V$ .<sup>4</sup>

$V_0$ <sup>5</sup>

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

<sup>6</sup>

$\rightarrow V_1$ <sup>7</sup>

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

<sup>8</sup>

Before check  $V$  get interrupted.<sup>9</sup>

$\rightarrow U_2$ <sup>10</sup>

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

<sup>11</sup>

# <sup>1</sup>*K-CAS* without timestamps

<sup>2</sup>Say our *K-CAS* solution just encapsulates every modifying operation {**Add**, **Remove**} into a *K-CAS* operation.

<sup>3</sup>Readers of the table may miss entries that are moved backwards by remove.

<sup>4</sup>Try find  $V$ .  
 $V_0$

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

<sup>5</sup>

<sup>6</sup> $\rightarrow V_1$

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------

<sup>7</sup>

<sup>8</sup>Before check  $V$  get interrupted.

<sup>9</sup> $\rightarrow U_2$

$\emptyset$	$X_0$	$Y_1$	$V_2$	$Z_2$	$W_2$
-------------	-------	-------	-------	-------	-------


<sup>10</sup>

<sup>11</sup>Delete  $Y$ . Move  $V$  back. Find  $Z$ , exit.  
 $V_2$

$\emptyset$	$X_0$	$V_1$	$Z_1$	$W_1$	$\emptyset$
-------------	-------	-------	-------	-------	-------------

<sup>12</sup>

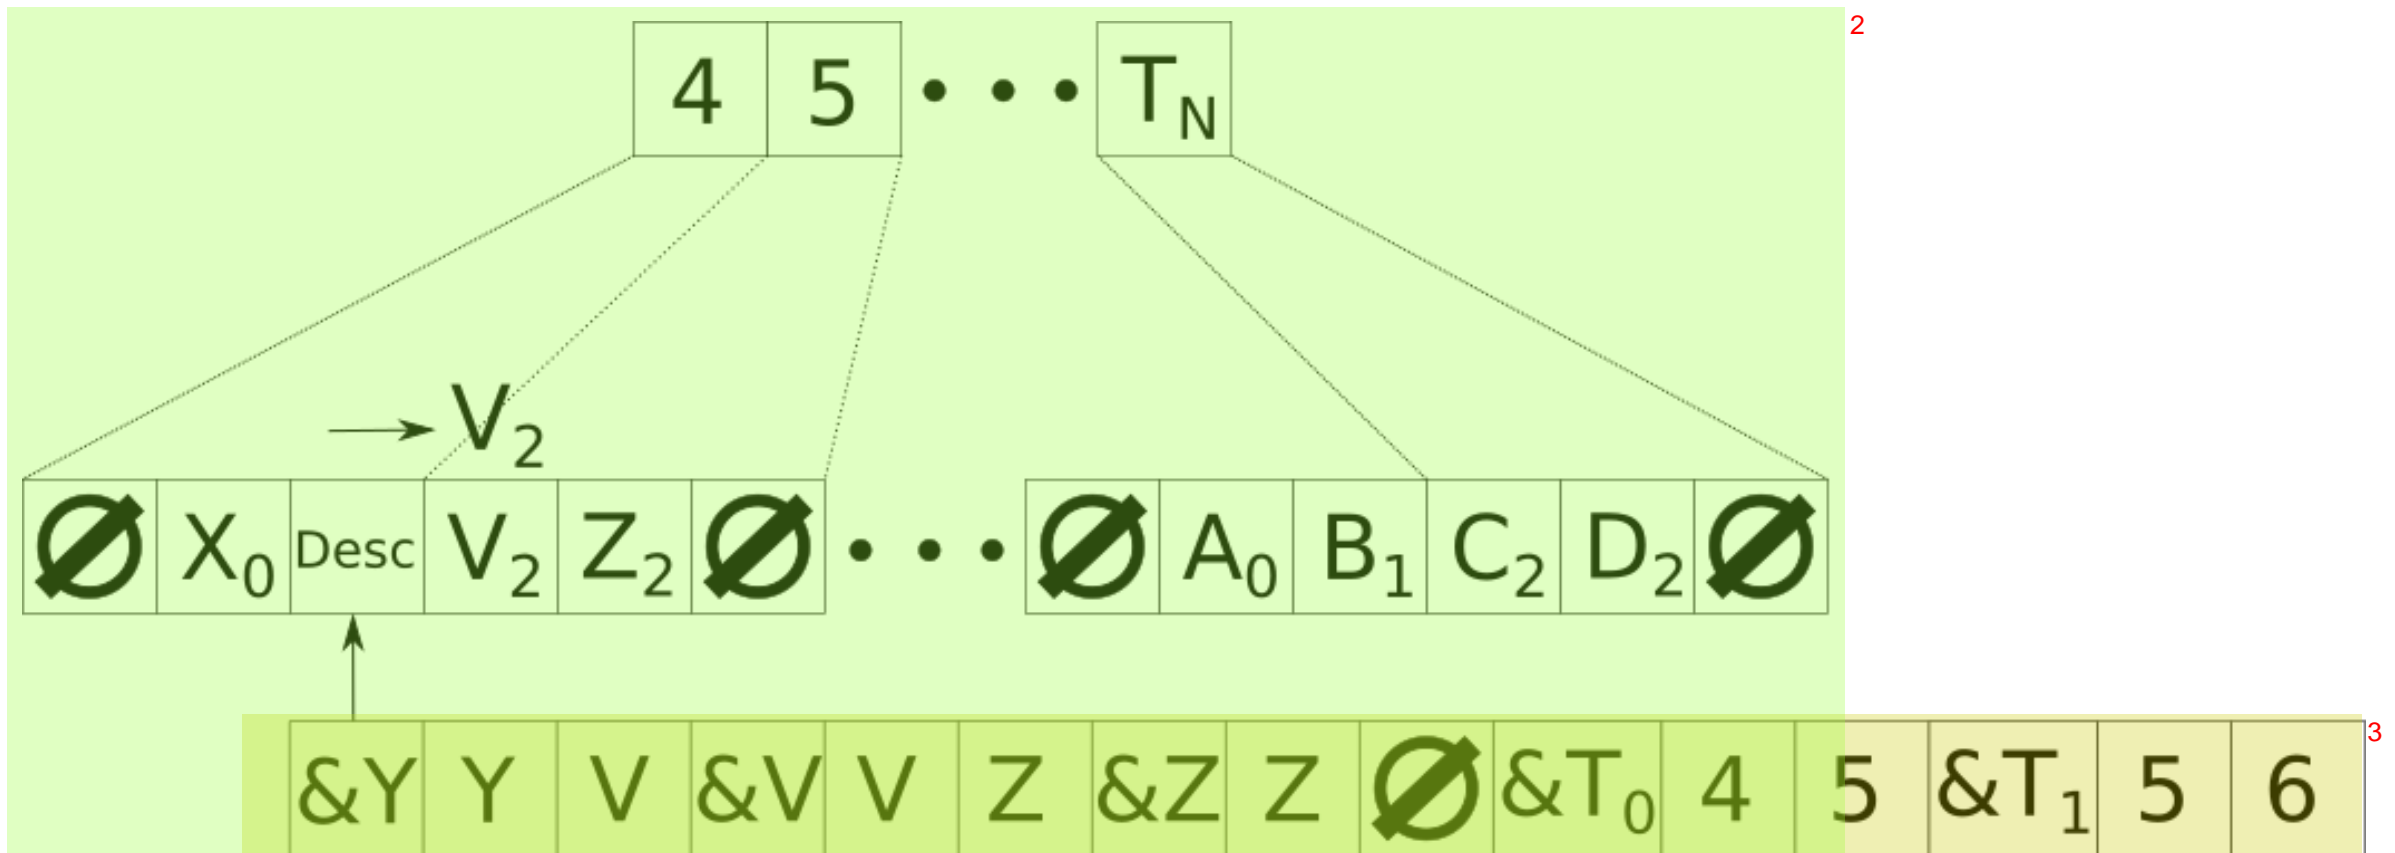
# Our Solution: Example<sup>1</sup>



Going to delete Y from table, with concurrent reader.<sup>3</sup>




# Our Solution: Example<sup>1</sup>



Ugly little array is a deletion descriptor.<sup>4</sup>

Moves items. Increments two timestamps.<sup>5</sup>

# Our Solution: Example<sup>1</sup>



Reader misses  $V$ , due to deletion of  $Y$ .<sup>3</sup>

Reader sees timestamp change, restarts operation.<sup>4</sup>

# Our Solution: Benefits<sup>1</sup>

# Our Solution: Benefits<sup>1</sup>

Relatively simple design, close to the sequential algorithm.<sup>2</sup>

# Our Solution: Benefits<sup>1</sup>

Relatively simple design, close to the sequential algorithm.<sup>2</sup>

No dynamic memory, great cache performance. Minimal memory overhead.<sup>3</sup>

Similar amount of CAS operations as bespoke dynamic memory solution.<sup>4</sup>

# Our Solution: Benefits<sup>1</sup>

Relatively simple design, close to the sequential algorithm.<sup>2</sup>

No dynamic memory, great cache performance. Minimal memory overhead.<sup>3</sup>

Similar amount of CAS operations as bespoke dynamic memory solution.<sup>4</sup>

Use of *K*-CAS allows for thread collaboration. Well defined non-blocking progress guarantees.<sup>5</sup>

Bulk relocation greatly reduces contention. Fast.<sup>6</sup>

# Our Solution: Correctness

# Our Solution: Correctness

Simple design means simple proof. Correctness is informally argued.



# Our Solution: Correctness

Simple design means simple proof. Correctness is informally argued.

Every modifying operation is a *K-CAS* operation. Cannot be seen midway.

Every reader must *remember* every timestamp seen.

Before any actions attempts to take effect they re-read timestamps. If any discrepancies are seen, retry operation.

# Our Solution: Progress

# Our Solution: Progress<sup>1</sup>

Solution is obstruction-free/lock-free. Obstruction-free **Contains**, lock-free **Add** and **Remove**.<sup>2</sup>

# Our Solution: Progress<sup>1</sup>

Solution is obstruction-free/lock-free. Obstruction-free **Contains**, lock-free **Add** and **Remove**.<sup>2</sup>

Every operation checks timestamps before the operation completes. Timestamps are coarse so operations can impede each other.<sup>3</sup>

The impeding of **Contains** means potentially no **Contains** will pass, but *at least* one **Add** or **Remove** will get through.<sup>4</sup>

# Benchmarking setup<sup>1</sup>

# Benchmarking setup<sup>1</sup>

## Hardware<sup>2</sup>

- 4 x Intel® Xeon® CPU E7-8890 v3, 18 cores each, 2 threads per core, 144 threads in total<sup>3</sup>
- HyperThreading avoided until the end
- PAPI used to measure various CPU artefacts
- *numactl* to control memory allocation

# Benchmarking setup<sup>1</sup>

## Hardware<sup>2</sup>

- 4 x Intel® Xeon® CPU E7-8890 v3, 18 cores each, 2 threads per core, 144 threads in total<sup>3</sup>
- HyperThreading avoided until the end
- PAPI used to measure various CPU artefacts
- *numactl* to control memory allocation

## Software<sup>4</sup>

- Microbenchmark measuring operations per microsecond<sup>5</sup>
- A number of strong performing concurrent hash tables
- Four load factors of 20%, 40%, 60%, and 80%
- Two read/write workloads of 10% and 20%

# Tables - Explainer<sup>1</sup>



# Tables - Explainer<sup>1</sup>

- Hopscotch Hashing [Herlihy, Shavit, Tzafrir; 2008]: Flattened separate chaining.<sup>2</sup>

# Tables - Explainer<sup>1</sup>

- Hopscotch Hashing [Herlihy, Shavit, Tzafrir; 2008]: Flattened separate chaining.<sup>2</sup>
- Lock-Free Linear Probing [Nielsen, Karlsson; 2016] : State simplified Purcell, Harris Table.

# Tables - Explainer<sup>1</sup>



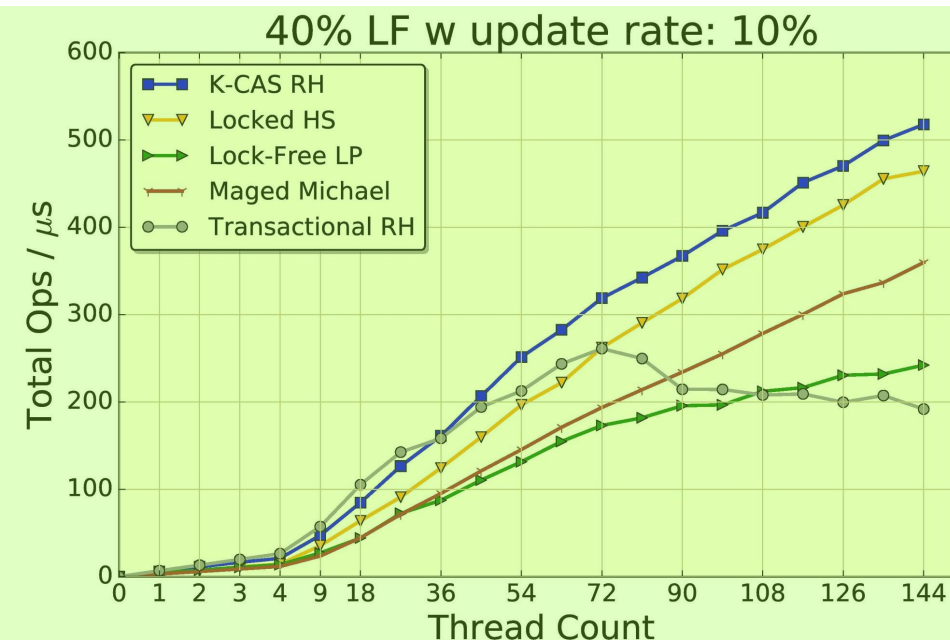

- Hopscotch Hashing [Herlihy, Shavit, Tzafrir; 2008]: Flattened separate chaining.
- Lock-Free Linear Probing [Nielsen, Karlsson; 2016] : State simplified Purcell, Harris Table.
- Separate Chaining [Maged Michael; 2003]: Per-bucket lock-free linked lists.

# Tables - Explainer<sup>1</sup>

- Hopscotch Hashing [Herlihy, Shavit, Tzafrir; 2008]: Flattened separate chaining.
- Lock-Free Linear Probing [Nielsen, Karlsson; 2016] : State simplified Purcell, Harris Table.
- Separate Chaining [Maged Michael; 2003]: Per-bucket lock-free linked lists.
- Lock-Elision Robin Hood. Serial algorithm with hardware transactional lock-elision wrapper
- *K-CAS* Robin Hood Hash. *K-CAS* with sharded timestamps.

# Performance 20%/40%

Total Operations per microsecond.







Number of threads.



# Performance 60%/80%

Total Operations per microsecond.



Number of threads.



# Performance Summary<sup>1</sup>

# Performance Summary<sup>1</sup>

Robin Hood scales best in almost all workloads.  
Otherwise very competitive with Hopscotch Hashing.<sup>2</sup>

Comfortably ahead (10%) with 10% update load.<sup>3</sup>  
More competitive (5%) with 20% updates.



# Performance Summary<sup>1</sup>

Robin Hood scales best in almost all workloads.  
Otherwise very competitive with Hopscotch Hashing.<sup>2</sup>

Comfortably ahead (10%) with 10% update load.<sup>3</sup>  
More competitive (5%) with 20% updates.

Robin Hood dominates other concurrent hash tables.<sup>4</sup>  
Gap narrows during Hyperthreading.

Transactional Robin Hood scales very strongly until  
Hyperthreading. Then it dies and never recovers.<sup>5</sup>

# Conclusion<sup>1</sup>

- First linearisable concurrent variant of Robin Hood Hashing.<sup>2</sup>
- Strong application of new *K*-CAS developments.
- Competitive performance compared to state of the art concurrent hash tables.

# Future Work<sup>3</sup>

- Extended Robin Hood work (different timestamp encodings/placements, cache aware, vectorised, various lock-based solutions)<sup>4</sup>
- Yahoo benchmark (YCSB)

**Thank you!**<sup>1</sup>

**Questions and  
Comments?**<sup>2</sup>

